

TASK 4 — ML Self-Tuner: Implementation and Test Cases (FINAL)

HACKSIMUL8 2026 · Department of Mechanical Engineering, PES University · 2 September 2026

Task 4 as stated: *"Implement the ML model and present test cases with suitable comparisons and merits of the utilised ML model."*

Status: COMPLETE. Mean ITAE improvement over the fixed PID: **+61.2%** across five test cases.

1. How the self-tuner works

Phase 1	0 → 3 s	fly on the fixed Task 2 gains, record the response
At 3 s		compute the 8-feature signature
		→ model → predicted Ki
		→ bound to the training range
		→ hand the integrator state over (bumpless transfer)
Phase 2	3 → 12 s	continue with the adapted gains

Three implementation details worth defending:

Detail	Why
Bumpless transfer	the integrator state carries across the switch, so the control signal does not jump when gains change
Bounded to the training range	the model may interpolate but never extrapolate — this is the fix for the failure described in §4
Only Ki is adapted	see §3 — it is the only gain the data reliably identifies

2. Results

Five test cases, deliberately chosen **outside the training grid**. Three controllers compared on each: the fixed Task 2 PID, our ML self-tuner, and an oracle whose gains are optimised directly for that exact condition.

Test case	ITAE fixed	ITAE ML	ITAE oracle	OS fixed	OS ML	Improvement
TC1 Heavy payload (m × 1.8)	18.762	3.627	0.301	0.00%	16.27%	+80.7%
TC2 Strong wind + gust	6.586	6.586	1.486	10.88%	10.88%	−0.0%
TC3 Large step + sustained tilt	5.803	1.541	0.302	0.00%	0.00%	+73.4%
TC4 Combined worst case	14.855	4.422	1.316	0.00%	6.24%	+70.2%
TC5 Sensor noise + payload	12.173	2.228	0.781	0.00%	2.54%	+81.7%
						mean +61.2%

Four of five cases improve by 70–82%. The gains the model chose:

Test case	Predicted Ki	Oracle Ki
TC1	11.43	12.23
TC2	0.00	19.52
TC3	2.36	3.11
TC4	8.03	8.39
TC5	6.84	8.39

On four of five, the predicted Ki lands within ~20% of the oracle's — from a 3-second observation, with no knowledge of the payload.

TC2 — why it shows exactly 0.0%, and why that is informative

The model predicted **Ki = 0.000**: it chose *not* to adapt, so the controller stayed on the baseline gains and the ITAE is identical.

This is not a silent failure. TC2 is the **lightest payload** case (mass × 1.1), so the thrust-ratio feature read close to nominal and the model correctly inferred that no payload correction was needed. What is actually hurting TC2 is **wind**, and the oracle's answer (Ki = 19.52) shows integral action would help — but wind is much less cleanly identifiable from the response signature than mass is.

The honest statement: the model adapts to what it can identify and declines to act when it cannot. That is preferable to a model that guesses confidently and gets it wrong — which is exactly what the earlier version did (§4).

3. Why only Ki is adapted

This is the central engineering decision in Task 4, and it comes directly from the Task 3 regression accuracy:

Gain	R ²	Adapted?
Ki	0.70 – 0.87	yes
Kp	0.21 – 0.52	no — held at the Task 2 value
Kd	0.14 – 0.37	no — held at the Task 2 value

Ki is uniquely determined. To cancel a steady-state error of a given size, integral action must have a specific strength. One right answer, so it is learnable.

Kp and Kd are not. Many (Kp, Kd) combinations give nearly identical cost — the objective has a flat valley in those directions — so `fminsearch` lands anywhere along it depending on its start point. **The training labels themselves are noisy**, and no model can predict noise.

The oracle confirms this is the right call: for TC1 it wants Kp = 20.2 against our baseline 30.2 (close), while Ki goes 0 → 12.2. **Ki is where the entire benefit lies.**

4. The failure that led here — and why it belongs in the report

The first working version let the model set **all three** gains. On unseen test conditions it extrapolated to **negative Kp**, the safety clamp floored it at 0.05, and the controller lost all proportional action:

BEFORE (all three gains adapted):

TC1	predicted Kp = 0.050	ITAE 152.046	vs fixed 18.762	-710.4%,	overshoot 353%
TC4	predicted Kp = 0.050	ITAE 31.074	vs fixed 14.855	-109.2%	
TC5	predicted Kp = 0.050	ITAE 29.577	vs fixed 12.173	-143.0%	
MEAN: -168.3% (i.e. far WORSE than doing nothing)					

AFTER (Ki only, bounded to the training range):

MEAN: +61.2%

Why this matters for the presentation: the diagnosis was not "tune the hyperparameters until it works". It was to ask *which parameters does my data actually identify?*, read the answer off the R^2 values, and restrict the model to those. That is a methodological fix, not a numerical one.

5. The oracle — and an honest reading of the gap

The oracle optimises gains directly for each test condition, so it bounds the best achievable. Mean remaining gap: **+455%**.

That number looks bad, and we should not hide it. Two things explain it:

1. **The oracle sees the full 12-second flight; the ML sees 3 seconds.** The oracle is allowed to optimise against information the self-tuner never has. It is a ceiling, not a fair competitor.
2. **The oracle adapts all three gains; we adapt one.** That is a deliberate restriction (§3), and it costs performance in exchange for reliability.

The meaningful comparison is against the **fixed PID**, which is what a real system would otherwise use, and there we improve by 61%. The oracle's role is to show how much headroom remains — and it shows there is real headroom, which is honest.

6. Merits

- **Genuinely self-tuning.** Never told the mass or the wind; infers conditions from its own response. Feature 6 exploits the fact that hover thrust equals weight.
- **Real-time capable.** Inference is one forward pass — microseconds. Re-running the optimiser online would take minutes and is impossible on a flight controller.
- **Substantial measured gain.** +61.2% mean ITAE improvement, +70–82% on four of five cases.
- **Fails safe.** Bounded to the training range; when the model cannot identify a condition it declines to adapt rather than guessing (TC2).
- **Trained entirely in simulation** — no flight testing, no hardware risk.
- **The restriction to K_i is derived from the data**, not assumed.

7. Limitations — state these

- **Envelope-bound.** Valid for mass $\times 1.0\text{--}2.4$ and wind ± 1 N. Outside that the model extrapolates; the bound prevents it from acting on such a prediction.
- **3-second identification window.** Until it has flown and measured, performance equals the baseline.
- **Wind is not well identified.** TC2 shows the model declining to adapt where the oracle would have. Adding a wind-specific feature — for instance the low-frequency component of the thrust residual — is the obvious next step.
- **K_p and K_d are not adapted at all**, so conditions requiring different transient shaping are not addressed.
- **Simulation-trained**, inheriting every modelling error: rigid airframe, no blade flapping, instantaneous motor response, no ground effect.
- **No formal stability guarantee.** A Lyapunov-based adaptive scheme would provide what a black-box regressor cannot.

8. Files

File	Role
<code>task4_test_ml_selftuning.m</code>	The Task 4 deliverable. Two-phase self-tuner, five test cases, three-way comparison with a multi-start oracle.
<code>task4_results.mat</code>	Per-test-case metrics, predicted and oracle gains
<code>task3_model.mat</code>	The trained model and its <code>predict_gains</code> handle
<code>sim_altitude.m</code>	Simulator — note <code>I0 / t0</code> for bumpless handover and <code>m_ctrl</code> for the unknown payload

How to run

```
cd 'C:\Users\LENOVO\Downloads\HACKSIMU8\solution'
task4_test_ml_selftuning           % ~9 min with parfor
```

Requires `task3_model.mat` and `task3_dataset.mat`. If either the cost function or the Task 2 baseline changes, re-run Task 3 first — otherwise Task 4 evaluates a model trained against a different problem.

9. Presenting Task 4 (about 3 minutes)

1. **"Fixed gains leave 36 cm of error with an unknown payload."** Show figure 05.
2. **"Our controller flies for 3 seconds, measures its own response, and adapts."** Walk the two-phase diagram.
3. **"It is never told the mass — it reads it off its own thrust command."** $r = +0.84$.
4. **"Here are five test cases outside the training grid."** Show the results table: +61.2% mean, +70–82% on four of five.
5. **"We adapt only K_i , because that is the only gain the data identifies."** Show the R^2 split and explain the flat cost valley.
6. **"We know this because the first version failed."** The –168% result, the diagnosis, the fix. This is a strength, not an admission.
7. **"And here is the headroom that remains."** The oracle gap, honestly framed.

Questions and answers

"Why does TC2 show zero improvement?" The model predicted $K_i = 0$ — it declined to adapt. TC2 is the lightest payload case, so the thrust-ratio feature read near nominal and the model correctly inferred no payload correction was needed. What hurts TC2 is wind, which is far less identifiable from the response signature. The model adapts to what it can identify and declines otherwise.

"Why not adapt all three gains?" We tried; it was 168% worse than doing nothing. K_p and K_d have R^2 around 0.2–0.5 because their labels are intrinsically noisy — many combinations give near-identical cost. K_i has R^2 0.70–0.87 because it is uniquely determined. Adapting only what the data identifies is the correct response.

"Your oracle gap is 455%. Isn't that bad?" The oracle optimises against the full 12-second flight and adapts all three gains; our self-tuner sees 3 seconds and adapts one. It is a ceiling, not a fair competitor. Against the fixed PID — what a real system would otherwise use — we improve 61%.

"How do you know it is safe?" Predictions are bounded to the range observed in training, the identification phase always runs on known-stable baseline gains, and the integrator handover is bumpless. There is no formal stability proof, which we state as a limitation.

"What would you do next?" Add a wind-specific feature to fix TC2, extend from altitude to all four channels, replace the fixed identification window with continuous adaptation, and add a Lyapunov-based stability guarantee.